

**A REPORT
ON
DETECTION OF FACE-MASK WORN BY A PERSON**

BY

HITESH ARYAN ACHARYA

2018AAPS0384H

AT

SMART i-ELECTRONICS SYSTEMS, PUNE

A Practice School-1 Station of

BIRLA INSTITUTE OF TECHNOLOGY AND SCIENCE, PILANI

June, 2020

**A REPORT
ON
DETECTION OF FACE-MASK WORN BY A PERSON**

BY

HITESH ARYAN ACHARYA

2018AAPS0384H

AT

SMART i-ELECTRONICS SYSTEMS, PUNE

A Practice School-1 Station of

BIRLA INSTITUTE OF TECHNOLOGY AND SCIENCE, PILANI

June, 2020

ACKNOWLEDGEMENTS

I am grateful to my Institute and the Practice School Division for organizing the Practice School-1 and setting up avenues for the same. I am deeply indebted to the Smart I - Electronics Systems for providing me with an opportunity to work with them and mentoring me. I would like to thank my instructor Prof. Rejesh N.A. for his constant guidance, clarification of my doubts and providing me with lecture videos and other invaluable learning material. I would also like to thank my mentor Mr. Pankaj Zanwar for explaining the details of the problem statement, helping me with the approach and providing me with the required material. Lat but not the least, I would like to thank my family and friends who have supported me throughout.

BIRLA INSTITUTE OF TECHNOLOGY AND SCIENCE PILANI
(RAJASTHAN)
Practice School Division

Station: Smart i-Electronics Systems

Centre: Pune

Duration: 6 weeks

Date of start: 18th May, 2020

Date of submission: 26th June, 2020

Title of the project: Detection of a Face-mask worn by a Person

ID No./Name(s)/ 2018AAPS0384H, Hitesh Aryan Acharya, B.E. ECE

Discipline(s)/
of the student(s)

Name(s) and Mr. Vaibhav, Employee
designation(s) Mr. Pankaj Zanwar, Director
of the expert(s):

Name(s) of Mr. Rejesh N. A.
the PS Faculty:

Key Words: Neural Networks, Propagation, Dataset, EPOCH, Optimization

Project Areas: Deep Learning, OpenCV

Abstract:

The following project involves building a neural network model which can identify and detect whether a person is wearing a face-mask using concepts like convolutions, data augmentation etc. It also requires understanding of TensorFlow/Keras and OpenCV.

Signature(s) of Student(s)

Signature of PS Faculty

Date

Date

Table of Contents

ACKNOWLEDGEMENTS	3
ABSTRACT SHEET	4
CHAPTER 1	7
Why is face-mask detection required?	7
Why use Deep-Learning?	7
CHAPTER 2	8
Convolutional Neural Networks	8
CHAPTER 3	9
3.1 Logistic Regression and Linear Regression	9
CHAPTER 4	11
4.1 IMPORTING THE NECESSARY LIBRARIES_	11
4.2 LOADING THE DATASET	12
4.3 UNDERSTANDING OUR DATASET	12
CHAPTER 5	14
General Architecture of the Neural Network	14
5.1 INPUT TO THE NEURAL NETWORK	15
CHAPTER 6	17
6.1 SIGMOID FUNCTION	17
6.2 INITIALIZING THE PARAMETERS	17
6.3 FORWARD & BACKWARD PROPAGATION	18
6.4 UPDATING THE PARAMETERS	20
CHAPTER 7	22
7.1 BINARY CLASSIFICATION	22
7.2 THE BINARY CLASSIFIER	23
7.3 SAMPLE TEST	24
7.4 COST	24
CHAPTER 8	26
8.1 Why we use TensorFlow and Keras	26
8.2 ImageDataGenerator	27
8.3 Cat Detection using TensorFlow	28
CHAPTER 9	30
9.1 Training the Neural Network for Face Mask Detection	30
9.2 Testing of our Model on Live Examples	36
9.3 Detection of Face Mask during a Live Video Stream	39

CHAPTER 10 **44**

 Loopholes in the Model 44

How can the Model be Improved 44

CONCLUSION **44**

APPENDIX-A..... **45**

APPENDIX-B **56**

APPENDIX-C **58**

APPENDIX-D **60**

APPENDIX-E **62**

APPENDIX-F **65**

REFERENCES..... **68**

GLOSSARY **69**

CHAPTER 1

Why is face-mask detection required?

Coronavirus as we all are aware is among the most contagious viruses. It is hence required that we all follow certain guidelines like wearing a face mask whenever outside. Social, religious and academic institutions, where there is a large influx of people, requires a robust technique to detect whether all are wearing a face-mask. Hence it is imperative that we develop a computer algorithm to do the same. This eliminates additional manpower and human error.

Why use Deep-Learning?

Deep learning uses multiple layers of neural networks to progressively extract higher level features from the raw input. For example, in image processing, lower layers may identify edges, while higher layers may identify the concepts relevant to a human such as digits or letters or faces.

Algorithms such as deep learning tend to perform better with more data. Living in such day and age where we have an abundance of data it is useful that we develop models which implement machine learning and deep learning techniques. The figure below shows how the performance of DL algorithms increase over amount of data compared to those of traditional learning algorithms.

Implementing a deep-learning algorithm to detect face-masks is very easy and effective. It can work on live-feed images and videos as well as can detect multiple face-masks at a time. It doesn't necessarily require a real dataset and yet still performs to the desired level. Most importantly, the code is very user friendly and can be tweaked to accommodate diverse parameters accordingly.




Figure 1- Deep Learning vs Traditional algorithms

CHAPTER 2

Convolutional Neural Networks

What are convolutions?

In image processing, convolutions are basically matrix operations on the pixel values to bring out a desired effect in the image. The name “convolutional neural network” indicates that the network employs a mathematical operation called convolution. Convolution is a specialized kind of linear operation. Convolutional networks are simply neural networks that use convolution in place of general matrix multiplication in at least one of their layers.

Why do we need convolutions?

Often images in the dataset aren't uniform and the characteristic properties of the image can be placed anywhere in the image. Hence we require a technique which can identify the important features of the image and make the image size uniform. Convolutions help condense the information present and reduce the size of the image at the same time. This pre-processing helps in reducing computation time. Pooling, Receptive Field, Weights etc. are some of the common terms in CNN which will be discussed later.



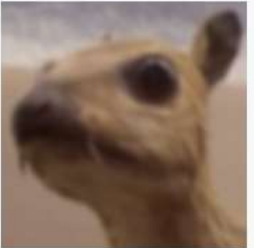
Operation	Kernel ω	Image result $g(x,y)$
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Gaussian blur 5×5 (approximation)	$\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$	

Figure 2- Effect of convolutions on image/Source- Wikipedia

CHAPTER 3

3.1 Logistic Regression and Linear Regression

Linear regression is a linear approach to modelling the relationship between a dependent variable and one or more independent variables. In **linear regression** we try to fit a straight line which passes through most of the data points. **Logistic Regression** is predominantly used in classification. It is a model used to model the probability of a certain event. But unlike **linear regression** where we compute the slope of a straight line, in **logistic regression** we have a S-shaped curve which performs binary classification, i.e. 0 or 1.

In **linear regression** we compute the slope of the straight line using the **least square fit** method. Here we measure the distance between the points and the straight line and add the squares of these distances. The minimum value pertaining to this is the best fit straight line.




Figure 3- Linear Regression

In **logistic regression** we try to divide the data into two broad classes using **maximum likelihood estimation**. In a nutshell we try to pass the majority of the data points through the S-shaped curve.




Figure 4- Logistic Regression

Such S-shaped curves are called logistic functions. Popular functions in machine learning include the [sigmoid function](#) and the [hyperbolic tan](#).

sigmoid :
$$A = \frac{1}{1+e^{-x}}$$

tanh(x) :
$$A = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$




Figure 5- Logistic vs Linear Regression

Here we will be learning to classify whether an image is a cat image or a non-cat image using **logistic regression**. Cat-image is represented by [1] in the labels' dataset and a non-cat image is represented by [0]. We will be using the [sigmoid function](#).

CHAPTER 4

LOGISTIC REGRESSION USING NEURAL NETWORK

4.1 IMPORTING THE NECESSARY LIBRARIES

```
import numpy as np
import matplotlib.pyplot as plt
import h5py
import scipy
from PIL import Image
from scipy import ndimage
from lr_utils import load_dataset

%matplotlib inline
```

Figure 6- Importing libraries

- **Numpy** is one of the most important library in Python. Using **numpy** we can compute array, vector and matrix operations with ease.
- **Matplotlib** is used for plotting graphs and images.
- **H5py** is important for binary data format.
- **Python Imaging Library** or **PIL** is a library that adds support for opening, manipulating, and saving many different image file formats.
- **SciPy** is library which is used to solve scientific and mathematical problems. It is built on the **NumPy** extension and allows the user to manipulate and visualize data with a wide range of high-level commands.
- **%matplotlib inline** sets the backend of matplotlib to the 'inline' backend. We use this in Jupyter notebook.

4.2 LOADING THE DATASET

```
# Loading the data (cat/non-cat)
train_set_x_orig, train_set_y, test_set_x_orig, test_set_y, classes = load_dataset()
```

Figure 7- Loading our dataset

We divide our datasets into two classes: the training set and testing set. These sets contain sub-classes with x as being the input and y being the output.

```
# Example of a picture
index = 20
plt.imshow(test_set_x_orig[index])
print ("y = " + str(train_set_y[:, index])+", it's a '"+classes[np.squeeze(train_set_y[:, index])].decode("utf-8") + "'picture.")
```

Figure 8- Example from the dataset

y = [1],it's a 'cat' picture.




Figure 98- It's a cat picture

This is an example from the training dataset. The output gives [1] if it is a cat picture.

4.3 UNDERSTANDING OUR DATASET

Colored images are 3-dimensional matrices. Pixel intensity across the x and y-axis represent two dimension while RGB makes up the third dimension. We use training examples to train our model.

```

m_train = train_set_x_orig.shape[0]
m_test = test_set_x_orig.shape[0]
num_px = train_set_x_orig.shape[1]

print ("Number of training examples: m_train = " + str(m_train))
print ("Number of testing examples: m_test = " + str(m_test))
print ("Height/Width of each image: num_px = " + str(num_px))
print ("Each image is of size: (" + str(num_px) + ", " + str(num_px) + ", 3)")
print ("train_set_x shape: " + str(train_set_x_orig.shape))
print ("train_set_y shape: " + str(train_set_y.shape))
print ("test_set_x shape: " + str(test_set_x_orig.shape))
print ("test_set_y shape: " + str(test_set_y.shape))

```

Figure 9- Extracting the shapes of the input images

```

Number of training examples: m_train = 209
Number of testing examples: m_test = 50
Height/Width of each image: num_px = 64
Each image is of size: (64, 64, 3)
train_set_x shape: (209, 64, 64, 3)
train_set_y shape: (1, 209)
test_set_x shape: (50, 64, 64, 3)
test_set_y shape: (1, 50)

```

Figure 10- Dimensions of the input images

We have 209 training and 50 testing examples. The number of pixels in the image is 64*64 with each axis having 64 pixels. The images form our input x. The output y is a row vector denoting [0] or [1] for each of the training example.

CHAPTER 5

General Architecture of the Neural Network

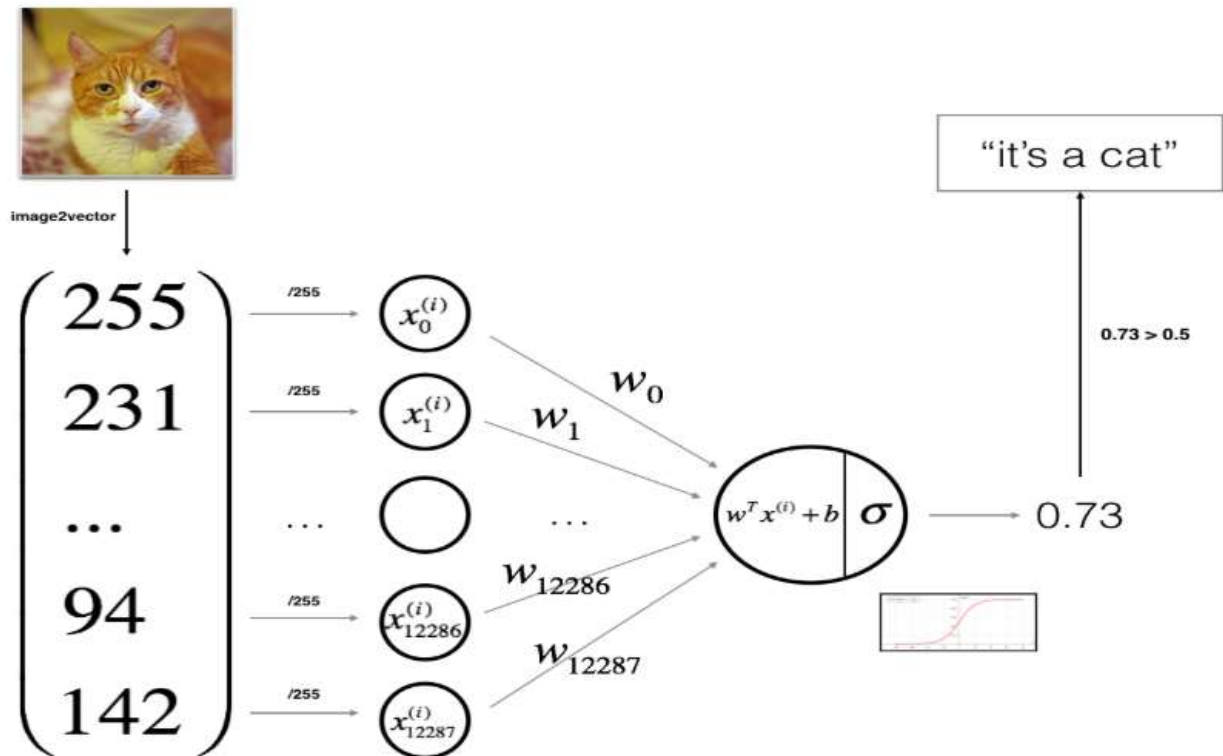


Figure 11- Model of the neural network/Source- Coursera-Introduction to Neural Networks and Deep Learning

Mathematical expression of the algorithm:

For one example $x^{(i)}$, forward propagation:

$$z^{(i)} = w^T \cdot x^{(i)} + b \quad (1)$$

$$\hat{y}^{(i)} = a^{(i)} = \text{sigmoid}(z^{(i)}) \quad (2)$$

$$L(a^{(i)}, y^{(i)}) = -y^{(i)} \log(a^{(i)}) - (1-y^{(i)}) \log(1-a^{(i)}) \quad (3)$$

The cost is then computed by summing over all training examples:

$$J = \frac{1}{m} \sum_{i=1}^m L(a^{(i)}, y^{(i)}) \quad (4)$$

Equations for back-propagation:

$$\frac{\partial J}{\partial w} = \frac{1}{m} X(A - Y)^T \quad (5)$$

$$\frac{\partial J}{\partial b} = \frac{1}{m} \sum_{i=1}^m (a^{(i)} - y^{(i)}) \quad (6)$$

As explained before, we first convert the 3D image into a unit dimensional column vector. We then normalize each row value by dividing it by 255. Each input feature is **matrix-multiplied** by a corresponding weight. We update these weights to get the desired result during back propagation. Since the weight vector has the same dimension as that of the input(**X_flatten**), we take the transpose of **w** vector and then perform matrix multiplication. We add a bias **b** to alter the weight according to our needs. The bias is a scalar parameter and also requires to be updated. The resultant is called **z**. The activation function is the **sigmoid function** which takes **z** as the input and outputs **a** i.e. the probability. Based on the probability we calculate the **loss function**. The **loss function** is based on the **maximum likelihood estimation**. It tells how much the predicted result is different from the actual result. The **cost function** is the mean of **loss function** over the entire training set.

Using the **cost function** we alter the weights in order to get the desired output. We achieve this by gradient descent which requires the computation of **dJ/dw** and **dJ/db**(denoted **dw** and **db** respectively) i.e. the change in L with respect to change in w and b. We go forward and backward till we get an acceptable low cost value.

NOTE :- The superscript (i) in the above equation denotes the i^{th} example. **A**, **Y** represent the vectorized version of **a**, **y** i.e. the **m** examples stacked together to get **m** columns.

5.1 INPUT TO THE NEURAL NETWORK

In order to input our images, we require to squeeze the 3 dimensions to a single dimension. This can be done by taking each pixel value in an orderly fashion and stacking it as a row vector. This way all pixel values are stacked one over the other making it a one-dimensional matrix. Further, each training example has to be computed independently. This can be done by taking an explicit for-loop and incrementing the loop to 50. But this is a very long process. To avoid such time consumption, we vectorize the input to a 2- dimensional input layer with each row representing the corresponding pixel of the input image and each column denoting the different examples. So the new dimensions of the input are now (64*64*3, 50).

A trick when you want to flatten a matrix X of shape (a,b,c,d) to a matrix **X_flatten** of shape (b*c*d, a) is to use:

```
X_flatten = X.reshape(X.shape[0], -1).T
```

```
# Reshape the training and test examples
train_set_x_flatten = train_set_x_orig.reshape(train_set_x_orig.shape[0], -1).T
test_set_x_flatten = test_set_x_orig.reshape(test_set_x_orig.shape[0], -1).T
print ("train_set_x_flatten shape: " + str(train_set_x_flatten.shape))
print ("train_set_y shape: " + str(train_set_y.shape))
print ("test_set_x_flatten shape: " + str(test_set_x_flatten.shape))
print ("test_set_y shape: " + str(test_set_y.shape))
print ("sanity check after reshaping: " + str(train_set_x_flatten[0:5,0]))
```

Figure 12- Reshaping the dimensions of our input

```
train_set_x_flatten shape: (12288, 209)
train_set_y shape: (1, 209)
test_set_x_flatten shape: (12288, 50)
test_set_y shape: (1, 50)
sanity check after reshaping: [17 31 56 22 33]
```

Figure 13- Dimensions after flattening

Each pixel value in our flattened image is ranging from 0-255. We normalize it by dividing each pixel by 255.

```
train_set_x = train_set_x_flatten/255.
test_set_x = test_set_x_flatten/255.
```

Figure 14- Normalizing the input

CHAPTER 6

Parameter Operations during Forward and Backward Propagations

6.1 SIGMOID FUNCTION

```
def sigmoid(z):
    s = 1/(1+np.exp(-z))
    return s
```

Figure 15- Defining the sigmoid function

6.2 INITIALIZING THE PARAMETERS

```
def initialize_with_zeros(dim):
    w = np.zeros([dim,1])
    b = 0

    assert(w.shape == (dim, 1))
    assert(isinstance(b, float) or isinstance(b, int))

    return w, b
```

Figure 16- Initializing the weights and bias

```
dim = 2
w, b = initialize_with_zeros(dim)
print ("w = " + str(w))
print ("b = " + str(b))
```

Figure 17- Calling the function

```
w = [[ 0.]
      [ 0.]]
b = 0
```

Figure 18-Initial value of the parameters

We initialize the weights and bias to zero. Since w is a matrix of the same dimension as the input, we multiply the transpose of w with X . The `assert` keyword is used to confirm the dimensions of the parameters. It gives an error if the dimensions are not matched. The bias b is arbitrary and hence we assign it with value 0.

6.3 FORWARD & BACKWARD PROPAGATION

FORWARD PROPAGATION

- The first layer consists of input features. The input features have to be a row vector.
- The second layer consists of two computations. They are pre-activation and activation respectively.
- Weights(w) are multiplied with the input features in the pre-activation. A bias(b) is also added to their product accordingly. It is denoted by z .
- The activation computes the probability of activation of the corresponding neuron. z is the input and the $g(z)$ is the output where $g(x)$ is the activation function.
- We often use the ReLU (Rectified Linear Unit) as the activation function in the hidden layers and Sigmoid function in the output layer.
- The maximum probability of activation in the output layer is usually the answer of our problem.

BACKWARD PROPAGATION

- During training, once the output activation is achieved, loss is calculated using the loss function. The loss is respect to the labelled data.
- We need to find the change in the loss function with respect to change in the weights and biases.
- We are essentially finding the minima of the loss function using a technique known as stochastic gradient descent.
- Stochastic gradient descent is a an optimizer algorithm which involves hyperparameters like the learning rate.

```
def propagate(w, b, X, Y):

    m = X.shape[1]

    A = sigmoid(np.dot(w.T, X) + b)
    cost = -np.sum((np.multiply(Y, np.log(A))) + (np.multiply(1-Y, np.log(1-A))))/m

    dw = np.dot(X, (A-Y).T)/m
    db = np.sum(A-Y)/m

    assert(dw.shape == w.shape)
    assert(db.dtype == float)
    cost = np.squeeze(cost)
    assert(cost.shape == ())

    grads = {"dw": dw,
             "db": db}

    return grads, cost
```

Figure 19- Computing cost during forward and backward propagation

```
w, b, X, Y = np.array([[1.],[2.]]), 2., np.array([[1.,2.,-1.],[3.,4.,-3.2]]), np.array([[1,0,1]])
grads, cost = propagate(w, b, X, Y)
print ("dw = " + str(grads["dw"]))
print ("db = " + str(grads["db"]))
print ("cost = " + str(cost))
```

Figure 20- Performing forward and backward propagation and storing the values in a dictionary

```
dw = [[ 0.99845601]
      [ 2.39507239]]
db = 0.00145557813678
cost = 5.80154531939
```

Figure 21- Change in cost function with respect to the weight and bias

For the forward propagation we calculate the **activation** and the **cost functions**. The equations have been defined above. 'm' represents the number of training examples. Since X is a 2 dimensional matrix, the columns denoting the examples hence **m = X.shape[1]**.

For the backward propagation the parameters **dw** and **db** are calculated. The equations have been defined above. Both w and **dw** have the same dimensions. Similarly for b and db. **np.squeeze ()** is a **numpy** function used to do away with redundant dimensions.

Finally we define a dictionary **grads** containing **dw** and **db**. The function returns **cost** and **grads**.

NOTE:- The above function takes in sample arguments for the sake of checking the correctness of the function. We will input the original arguments shortly.

6.4 UPDATING THE PARAMETERS

```
def optimize(w, b, X, Y, num_iterations, learning_rate, print_cost = False):
    costs = []

    for i in range(num_iterations):
        grads, cost = propagate(w, b, X, Y)

        dw = grads["dw"]
        db = grads["db"]

        w -= learning_rate*dw
        b -= learning_rate*db

        if i % 100 == 0:
            costs.append(cost)

        if print_cost and i % 100 == 0:
            print ("Cost after iteration %i: %f" %(i, cost))

    params = {"w": w,
              "b": b}

    grads = {"dw": dw,
             "db": db}

    return params, grads, costs
```

Figure 22- Updating the weight and bias

```
params, grads, costs = optimize(w, b, X, Y, num_iterations= 100, learning_rate = 0.009, print_cost = False)

print ("w = " + str(params["w"]))
print ("b = " + str(params["b"]))
print ("dw = " + str(grads["dw"]))
print ("db = " + str(grads["db"]))
```

Figure 23- Extracting the value of parameters after updating

```
w = [[ 0.19033591]
      [ 0.12259159]]
b = 1.92535983008
dw = [[ 0.67752042]
       [ 1.41625495]]
db = 0.219194504541
```

Figure 24- Value of parameters

Here we are updating the parameters using gradient descent. First we retrieve the parameters `dw` and `db` from the `grads` dictionary. Using the hyperparameter `learning_rate` = 0.009, we compute `w` and `b` using the above equations. We store the updated parameters in dictionaries `params` and `grads`. The function returns `params`, `grads` and `costs`.

We have additionally defined a `costs` array. The `cost` gets updated after every iteration. We append the value of the `cost` variable in the costs array after every 100 iterations. Using this array we will plot the learning curve.

NOTE:- The above function takes in sample arguments for the sake of checking the correctness of the function. We will input the original arguments shortly.

CHAPTER 7

7.1 BINARY CLASSIFICATION

```
def predict(w, b, X):

    m = X.shape[1]
    Y_prediction = np.zeros((1,m))
    w = w.reshape(X.shape[0], 1)

    A = sigmoid(np.dot(w.T, X) + b)

    for i in range(A.shape[1]):

        if(A[0][i] <= 0.5):
            Y_prediction[0][i] = 0

        else:
            Y_prediction[0][i] = 1
        pass

    assert(Y_prediction.shape == (1, m))

    return Y_prediction
```

Figure 25- Predicting an example

```
w = np.array([[0.1124579],[0.23106775]])
b = -0.3
X = np.array([[1.,-1.1,-3.2],[1.2,2.,0.1]])
print ("predictions = " + str(predict(w, b, X)))
```

Figure 26- Testing an example on the predict function

```
predictions = [[ 1.  1.  0.]
```

Figure 27- Output for our test example

Here we classify the image as a cat image or a non-cat image based on the probability computed by the activation function. If the probability is greater than 0.5 then it is classified as a cat image and a non-cat image otherwise. The dimensions of **Y_prediction** is the same as **A**, i.e. (1, m). It classifies each example as a cat or a non-cat image.

NOTE:- The above function takes in sample arguments for the sake of checking the correctness of the function. We will input the original arguments shortly.

7.2 THE BINARY CLASSIFIER

```
def model(X_train, Y_train, X_test, Y_test, num_iterations = 2000, learning_rate = 0.5, print_cost = False):
    w, b = np.zeros([X_train.shape[0], 1]), 0

    parameters, grads, costs = optimize(w, b, X_train, Y_train, num_iterations, learning_rate, print_cost = False)

    w = parameters["w"]
    b = parameters["b"]

    Y_prediction_test = predict(w, b, X_test)
    Y_prediction_train = predict(w, b, X_train)

    print("train accuracy: {} %".format(100 - np.mean(np.abs(Y_prediction_train - Y_train)) * 100))
    print("test accuracy: {} %".format(100 - np.mean(np.abs(Y_prediction_test - Y_test)) * 100))

    d = {"costs": costs,
        "Y_prediction_test": Y_prediction_test,
        "Y_prediction_train": Y_prediction_train,
        "w": w,
        "b": b,
        "learning_rate": learning_rate,
        "num_iterations": num_iterations}

    return d
```

Figure 28- Linking all the functions and making a final model

```
d = model(train_set_x, train_set_y, test_set_x, test_set_y, num_iterations = 2000, learning_rate = 0.005, print_cost = True)
```

Figure 29- Setting the hyperparameters and providing the training and testing dataset

```
train accuracy: 99.04306220095694 %
test accuracy: 70.0 %
```

Figure 30- Displaying the accuracy

Here we build the neural network that we intended from the beginning. We pass the original parameters. All the functions defined before are called through cascading. Finally we test the accuracy of our model by comparing **Y_prediction** with **train_set_y** and **test_set_y**. We see the training accuracy to be much greater than the testing accuracy. This is a case of overfitting. But the model performs pretty good considering the network uses no hidden layers.

7.3 SAMPLE TEST

```
index = 49
plt.imshow(test_set_x[:,index].reshape((num_px, num_px, 3)))
print("y = "+str(test_set_y[0,index])+", you predicted that it is a \""+classes[d["y_prediction_test"][0,index]].decode("utf-8")
```

Figure 31- Displaying a result from testing

y = 0, you predicted that it is a "non-cat" picture.




Figure 32- The desired result

This is a test example. We see the neural network predicts the desired result.

7.4 COST

```
# Plot learning curve (with costs)
costs = np.squeeze(d['costs'])
plt.plot(costs)
plt.ylabel('cost')
plt.xlabel('iterations (per hundreds)')
plt.title("Learning rate = " + str(d["learning_rate"]))
plt.show()
```

Figure 33- Plot learning curve




Figure 34- The cost function vs no. of iterations with a learning rate = 0.005

Here we plot the cost as a function of number of iterations. Here we have hard coded the learning rate. We see the cost decreases with increase in the number of iterations.

CHAPTER 8

8.1 Why we use Tensorflow and Keras

TensorFlow is a library package in python which is a symbolic math library and has extensive use in the fields of data science and machine learning. The in-built functions make the code reusable and efficient.

Keras is an open-source neural network library written in python which is used on top of **TensorFlow**. It focuses on being user-friendly, modular and extensible.

- In the following program we use pooling alongside convolutions. Pooling further reduces the size of the image.
- In MaxPooling it retains the maximum value of the specified neighborhood and discards the other pixel values.
- In AveragePooling it averages the values in the specified neighborhood.




Figure 35- Max Pooling




Figure 36- Average Pooling

8.2 ImageDataGenerator

Another feature in [TensorFlow](#) is its ability to label the datasets based on the sub-directories. Most of the times, the dataset generated are not labelled and it is a very time-consuming activity to label each and every image. In [TensorFlow](#) there is already a in-built function **ImageDataGenerator** which label the datasets.

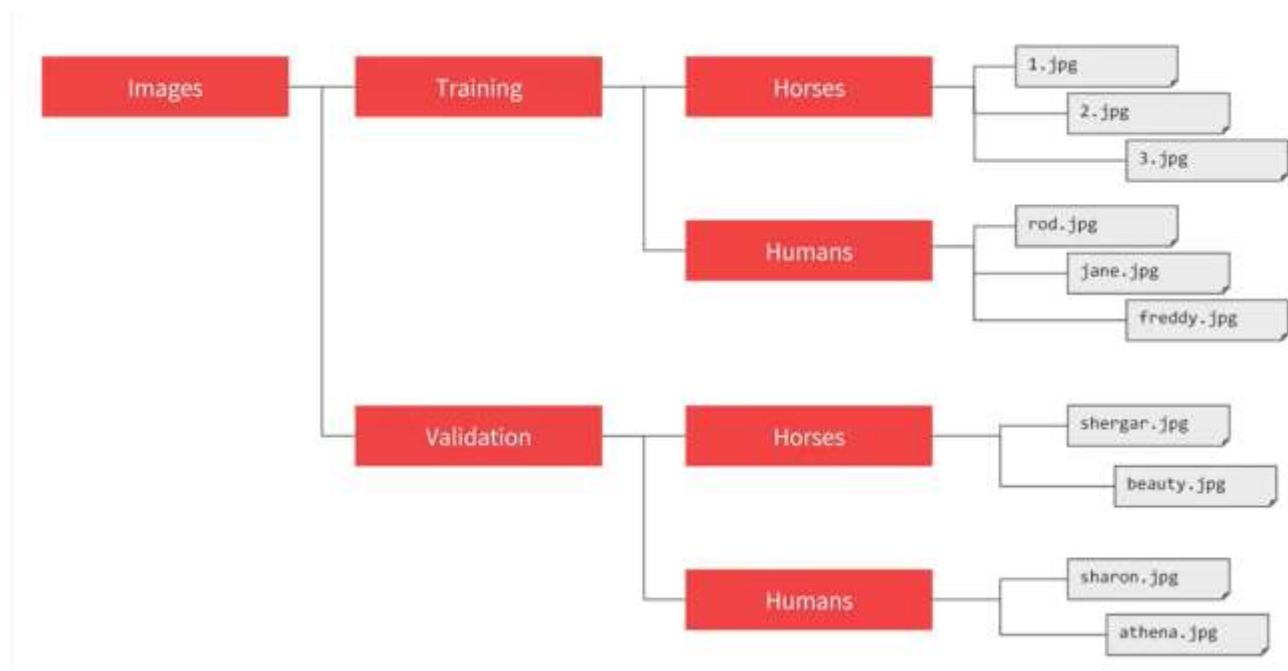


Figure 37- Coursera, Introduction to TensorFlow for Artificial Intelligence, Machine Learning, and Deep Learning

8.3 Cat Detection using TensorFlow

Here I train a similar neural network used for detecting images of cats using [TensorFlow](#) and [Keras](#).

```
import tensorflow as tf
import os
import zipfile
from os import path, getcwd, chdir

path = f"{getcwd()}/../tmp2/happy-or-sad.zip"

zip_ref = zipfile.ZipFile(path, 'r')
zip_ref.extractall("/tmp/h-or-s")
zip_ref.close()
```

Figure 38- Importing zip files and extracting them

```
def train_cat_not_model():

    DESIRED_ACCURACY = 0.999

    class myCallback(tf.keras.callbacks.Callback):

        def on_epoch_end(self, epoch, logs = {}):

            if(logs.get('acc') > DESIRED_ACCURACY):
                print("Reached 99.9% accuracy so cancelling training!")
                self.model.stop_training = True

    callbacks = myCallback()
```

Figure 39- This function trains the model till it achieves the desired accuracy

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3,3), activation = 'relu', input_shape = (150, 150, 3)),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Conv2D(32, (3,3), activation = 'relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Conv2D(32, (3,3), activation = 'relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation = 'relu'),
    tf.keras.layers.Dense(1, activation = 'sigmoid')
])
```

Figure 40- Applying convolutions and pooling to reduce computation time. The 'dense' keyword defines the number of neurons in the layer.

```

from tensorflow.keras.optimizers import RMSprop

model.compile(loss = 'binary_crossentropy', optimizer = RMSprop(lr = 0.001), metrics = ['accuracy'])

from tensorflow.keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(rescale = 1./255)

train_generator = train_datagen.flow_from_directory(
    '/tmp/h-or-s',
    target_size = (150, 150),
    batch_size = 10,
    class_mode = 'binary')

```

Figure 41- Propagating through the neural network and generating labels for the dataset. The rescale normalizes the pixel values

```

history = model.fit_generator(
    train_generator, steps_per_epoch = 8, epochs = 15, verbose = 1, callbacks = [callbacks])

return history.history['acc'][-1]

```

Figure 42- Training the model till the desired accuracy

`train_cat_not_model()`

Figure 43- Calling the function

OUTPUT:

```

Found 80 images belonging to 2 classes.
Epoch 1/15
8/8 [=====] - 5s 638ms/step - loss: 1.0556 - acc: 0.5875
Epoch 2/15
8/8 [=====] - 1s 71ms/step - loss: 0.7459 - acc: 0.7625
Epoch 3/15
8/8 [=====] - 1s 64ms/step - loss: 0.2174 - acc: 0.9125
Epoch 4/15
8/8 [=====] - 1s 73ms/step - loss: 0.1534 - acc: 0.9500
Epoch 5/15
8/8 [=====] - 1s 63ms/step - loss: 0.1707 - acc: 0.9250
Epoch 6/15
8/8 [=====] - 1s 63ms/step - loss: 0.0388 - acc: 0.9875
Epoch 7/15
7/8 [=====>....] - ETA: 0s - loss: 0.0333 - acc: 1.0000Reached 99.9% accuracy so cancelling training!
8/8 [=====] - 1s 73ms/step - loss: 0.0376 - acc: 1.0000

```

Figure 44- Training till the model achieves 99.9% accuracy, i.e., till the 7th epoch

CHAPTER 9

9.1 Training the Neural Network for Face Mask Detection

1) Importing Libraries and Packages

```

2      # python train_mask_detector.py --dataset dataset
3
4      # import the necessary packages
5      from tensorflow.keras.preprocessing.image import ImageDataGenerator
6      from tensorflow.keras.applications import MobileNetV2
7      from tensorflow.keras.layers import AveragePooling2D
8      from tensorflow.keras.layers import Dropout
9      from tensorflow.keras.layers import Flatten
10     from tensorflow.keras.layers import Dense
11     from tensorflow.keras.layers import Input
12     from tensorflow.keras.models import Model
13     from tensorflow.keras.optimizers import Adam

```

- We import all the necessary dependencies from TensorFlow and Keras. Some of the functions have already been discussed.
- ImageDataGenerator is used in data augmentation and class labelling.
- MobileNetV2 is a neural network architecture optimised for mobile devices.
- AveragePooling, Dropout, Flatten Dense, Input are the usual functions that are required to build a neural network.
- We'll be using Adam optimizer in our deep net.

```

14     from tensorflow.keras.applications.mobilenet_v2 import preprocess_input
15     from tensorflow.keras.preprocessing.image import img_to_array
16     from tensorflow.keras.preprocessing.image import load_img
17     from tensorflow.keras.utils import to_categorical
18     from sklearn.preprocessing import LabelBinarizer
19     from sklearn.model_selection import train_test_split
20     from sklearn.metrics import classification_report
21     from imutils import paths
22     import matplotlib.pyplot as plt
23     import numpy as np
24     import argparse
25     import os

```

- Pre-processing the input is required in neural networks especially when using a system with low computation power.
- Img_to_array is used to convert the multi-dimensional input image to a single dimension.
- To_categorical and LabelBinarizer are ways to label an output. It makes the training easier and efficient. These functions have been explained in Appendix-D.
- Sklearn is an additional machine learning library. It features various classification, regression and clustering algorithms including support vector machines etc.

The complete code without explanation is available in Appendix-A.

- Imutils is an important convenience package. Paths is used to store the path of the dataset and other testing examples for the model to use.
- Matplotlib, NumPy, argparse and os have additional utilities as we'll see later.

2) Argparsing Arguments

```

27  # construct the argument parser and parse the arguments
28  ap = argparse.ArgumentParser()
29  ap.add_argument("-d", "--dataset", required=True,
30                  help="path to input dataset")
31  ap.add_argument("-p", "--plot", type=str, default="plot.png",
32                  help="path to output loss/accuracy plot")
33  ap.add_argument("-m", "--model", type=str,
34                  default="mask_detector.model",
35                  help="path to output face mask detector model")
36  args = vars(ap.parse_args())

```

- Argparsing has been discussed in detail in Appendix A.
- Dataset argument is required to train the images. Dataset contains all our training and testing examples. This is a necessary argument.
- The plot displays the training and validation loss and accuracy over number of iterations.
- The neural network model is stored in mask_detector.model as default. If necessary we can change the default parameter.
- We store the parse arguments in a dictionary args. This has been discussed in Appendix-A.

3) Hyperparameter Initialization and Loading Dataset

```

38  # initialize the initial learning rate, number of epochs to train for,
39  # and batch size
40  INIT_LR = 1e-4
41  EPOCHS = 20
42  BS = 32

```

- We store the values of the hyper-parameters. Note that the learning rate is hard-coded for the time being. As we will see later, we add a decay to the learning rate. BS denotes Batch-size.

```

44  # grab the list of images in our dataset directory, then initialize
45  # the list of data (i.e., images) and class images
46  print("[INFO] loading images...")
47  imagePath = list(paths.list_images(args["dataset"]))
48  data = []
49  labels = []

```

- We create a list containing all the dataset images. We use the function paths.list_images to get the path of our dataset. Args is the dictionary we defined during argparsing.
- We create two lists – one for the data and the other for the labels.

The complete code without explanation is available in Appendix-A.


```

51 # Loop over the image paths
52 for imagePath in imagePaths:
53     # extract the class label from the filename
54     label = imagePath.split(os.path.sep)[-2]
55
56     # load the input image (224x224) and preprocess it
57     image = load_img(imagePath, target_size=(224, 224))
58     image = img_to_array(image)
59     image = preprocess_input(image)
60
61     # update the data and labels lists, respectively
62     data.append(image)
63     labels.append(label)
64
65 # convert the data and labels to NumPy arrays
66 data = np.array(data, dtype="float32")
67 labels = np.array(labels)

```

- We loop over the imagePaths list and extract the class label from the filename. imagePath.split is a list containing the names of the directories that make up the path. The second last element of the list is the name of the labels' file.
- We perform some operations on the input image. This includes converting the 3D image to a single dimensional array and pre-processing it. Pre-processing has been discussed in Appendix-C.
- The lists – label and data are being appended after each iteration of the loop.
- We convert these lists to NumPy arrays for the ease of operations.

4) Processing the Dataset

```

69 # perform one-hot encoding on the labels
70 lb = LabelBinarizer()
71 labels = lb.fit_transform(labels)
72 labels = to_categorical(labels)
73

```

- LabelBinarizer, fit_transform and to_categorical have been discussed in Appendix-D.

```

74 # partition the data into training and testing splits using 75% of
75 # the data for training and the remaining 25% for testing
76 (trainX, testX, trainY, testY) = train_test_split(data, labels,
77     test_size=0.20, stratify=labels, random_state=42)
78

```

- We split our dataset using the train_test_split function. It splits the data and labels according to the proportion assigned to test_size. Stratify ensures that the training and testing samples contain an equal measure of dichotomy.
- Random_state = 42 ensures that the division between test and train samples are the same every time the program is run. Random_state = 0 would make the partition random at every run.

5) Data Augmentation

```

79  # construct the training image generator for data augmentation
80  aug = ImageDataGenerator(
81      rotation_range=20,
82      zoom_range=0.15,
83      width_shift_range=0.2,
84      height_shift_range=0.2,
85      shear_range=0.15,
86      horizontal_flip=True,
87      fill_mode="nearest")

```

- We create an instance of data augmentation using the ImageDataGenerator class. We will apply this to our dataset later.

6) Building the Neural Network

```

89  # Load the MobileNetV2 network, ensuring the head FC layer sets are
90  # left off
91  baseModel = MobileNetV2(weights="imagenet", include_top=False,
92      input_tensor=Input(shape=(224, 224, 3)))

```

- We create a base model which will be the input layer for our neural network. The function arguments have their usual meanings. MobileNetV2 architecture has been discussed in detail in Appendix-D.

```

94  # construct the head of the model that will be placed on top of the
95  # the base model
96  headModel = baseModel.output
97  headModel = AveragePooling2D(pool_size=(7, 7))(headModel)
98  headModel = Flatten(name="flatten")(headModel)
99  headModel = Dense(128, activation="relu")(headModel)
100 headModel = Dropout(0.5)(headModel)
101 headModel = Dense(2, activation="softmax")(headModel)

```

- Further layers of the neural network are defined.
- These layers are placed on top of the base model, i.e., the output from the base model will be loaded into the head model.
- We apply average pooling with (7,7) blocks. This is similar to max pooling with the difference being we take the average of the (7,7) block instead of taking the max value.
- We define a layer with 128 units and relu activation function.
- We apply dropout with the threshold value being 0.5.
- The output layer has 2 units – with mask or without mask classes. The output activation is the softmax function.

```

102
103 # place the head FC model on top of the base model (this will become
104 # the actual model we will train)
105 model = Model(inputs=baseModel.input, outputs=headModel)
106
107 # Loop over all layers in the base model and freeze them so they will
108 # *not* be updated during the first training process
109 for layer in baseModel.layers:
110     layer.trainable = False

```

- We construct the model using Model function. Note that the output of the base model is the input to the head model.
- We freeze the layers of the base model. This ensures that the parameters are not updated during the first iteration of training.

```

112 # compile our model
113 print("[INFO] compiling model...")
114 opt = Adam(lr=INIT_LR, decay=INIT_LR / EPOCHS)
115 model.compile(loss="binary_crossentropy", optimizer=opt,
116               metrics=["accuracy"])
117

```

- We compile the model and define the hyper-parameters. The decay is also defined which is essentially reducing the value of the learning rate.
- The loss is computed using binary cross entropy.

7) Training the Neural Network

```

118 # train the head of the network
119 print("[INFO] training head...")
120 H = model.fit(
121     aug.flow(trainX, trainY, batch_size=BS),
122     steps_per_epoch=len(trainX) // BS,
123     validation_data=(testX, testY),
124     validation_steps=len(testX) // BS,
125     epochs=EPOCHS)

```

- Now we train the neural network using our training dataset.
- Our validation dataset is same as the testing dataset. It validates the test dataset after each propagation and compares it with test labels.

8) Validation and Saving the Model

```

127     # make predictions on the testing set
128     print("[INFO] evaluating network...")
129     predIdxs = model.predict(testX, batch_size=BS)
130
131     # for each image in the testing set we need to find the index of the
132     # label with corresponding largest predicted probability
133     predIdxs = np.argmax(predIdxs, axis=1)
134
135     # show a nicely formatted classification report
136     print(classification_report(testY.argmax(axis=1), predIdxs,
137                               target_names=lb.classes_))

```

- We test our trained model and store the values in a variable `predIdxs`.
- We use `argmax` function to get the predicted results. `Argmax` gives out the input corresponding to the max value of the function. In this case, if the maximum probability corresponds to an image with mask, the `argmax` function will give the output – `with_mask`.

```

139     # serialize the model to disk
140     print("[INFO] saving mask detector model...")
141     model.save(args["model"], save_format="h5")
142

```

- We save the model in the h5 format.

9) Plotting the different Parameters

```

143     # plot the training loss and accuracy
144     N = EPOCHS
145     plt.style.use("ggplot")
146     plt.figure()
147     plt.plot(np.arange(0, N), H.history["loss"], label="train_loss")
148     plt.plot(np.arange(0, N), H.history["val_loss"], label="val_loss")
149     plt.plot(np.arange(0, N), H.history["accuracy"], label="train_acc")
150     plt.plot(np.arange(0, N), H.history["val_accuracy"], label="val_acc")
151     plt.title("Training Loss and Accuracy")
152     plt.xlabel("Epoch #")
153     plt.ylabel("Loss/Accuracy")
154     plt.legend(loc="lower left")
155     plt.savefig(args["plot"])

```

- We plot the training accuracy, training loss, validation accuracy and validation loss over iterations. Finally we save the plot with `.png` extension.

9.2 Testing of our Model on Live Examples

1) Importing the Libraries

```

2  # python detect_mask_image.py --image examples/example_01.png
3
4  # import the necessary packages
5  from tensorflow.keras.applications.mobilenet_v2 import preprocess_input
6  from tensorflow.keras.preprocessing.image import img_to_array
7  from tensorflow.keras.models import load_model
8  import numpy as np
9  import argparse
10 import cv2
11 import os
12

```

- Import the necessary packages and dependencies.
- cv2 is the OpenCV library and is necessary for face detections on live examples.

2) Argparsing Arguments

```

13 # construct the argument parser and parse the arguments
14 ap = argparse.ArgumentParser()
15 ap.add_argument("-i", "--image", required=True,
16                 help="path to input image")
17 ap.add_argument("-f", "--face", type=str,
18                 default="face_detector",
19                 help="path to face detector model directory")
20 ap.add_argument("-m", "--model", type=str,
21                 default="mask_detector.model",
22                 help="path to trained face mask detector model")
23 ap.add_argument("-c", "--confidence", type=float, default=0.5,
24                 help="minimum probability to filter weak detections")
25 args = vars(ap.parse_args())

```

- We parse arguments. The image argument is necessary to parse. The model will try to detect the face mask on this image. We give the path to the image.
- We create other default arguments. These will be used in our program.
- We have discussed about confidence and confidence threshold in detail. The confidence here is 0.5.
- We create a dictionary args.

3) Loading the Face Detection Model

```

27 # Load our serialized face detector model from disk
28 print("[INFO] loading face detector model...")
29 prototxtPath = os.path.sep.join([args["face"], "deploy.prototxt"])
30 weightsPath = os.path.sep.join([args["face"],
31                                 "res10_300x300_ssd_iter_140000.caffemodel"])
32 net = cv2.dnn.readNet(prototxtPath, weightsPath)
33

```

- This is a built in function in OpenCV to detect faces in an image. Args["face"] returns face_detector since during argparsing we defined its default value the same.

The complete code without explanation is available in Appendix-A.

- The net variable contains the information on how to detect a face in an image.

```

34 # Load the face mask detector model from disk
35 print("[INFO] loading face mask detector model...")
36 model = load_model(args["model"])

```

- We load the trained model from the disk. Note that model was created during argparsing.

4) Face Detection in the Input Image

```

38 # Load the input image from disk, clone it, and grab the image spatial
39 # dimensions
40 image = cv2.imread(args["image"])
41 orig = image.copy()
42 (h, w) = image.shape[:2]
43
44 # construct a blob from the image
45 blob = cv2.dnn.blobFromImage(image, 1.0, (300, 300),
46                               (104.0, 177.0, 123.0))
47
48 # pass the blob through the network and obtain the face detections
49 print("[INFO] computing face detections...")
50 net.setInput(blob)
51 detections = net.forward()

```

- We create an instance of the image, clone it and grab the spatial dimensions of the image.
- We construct a blob of the image. The function blobFromImage has been discussed in detail in Appendix-C.
- We input the blob to our face detection model which is net. It then detects the face and the information is then stored in another variable detections.

5) Building a Confidence Threshold for Face Detection

```

53 # Loop over the detections
54 for i in range(0, detections.shape[2]):
55     # extract the confidence (i.e., probability) associated with
56     # the detection
57     confidence = detections[0, 0, i, 2]
58
59     # filter out weak detections by ensuring the confidence is
60     # greater than the minimum confidence
61     if confidence > args["confidence"]:
62         # compute the (x, y)-coordinates of the bounding box for
63         # the object
64         box = detections[0, 0, i, 3:7] * np.array([w, h, w, h])
65         (startX, startY, endX, endY) = box.astype("int")

```

- Only images where face detection is above a certain confidence threshold are selected.
- After they pass the confidence threshold, we grab the spatial dimensions of the face. We ensure dimensions are integer type by using astype("int").

```

67     # ensure the bounding boxes fall within the dimensions of
68     # the frame
69     (startX, startY) = (max(0, startX), max(0, startY))
70     (endX, endY) = (min(w - 1, endX), min(h - 1, endY))

```

- Here we check the validity of the box dimensions by ensuring they are inside the image spatial dimensions.

6) Pre-processing the Input Image

```

72     # extract the face ROI, convert it from BGR to RGB channel
73     # ordering, resize it to 224x224, and preprocess it
74     face = image[startY:endY, startX:endX]
75     face = cv2.cvtColor(face, cv2.COLOR_BGR2RGB)
76     face = cv2.resize(face, (224, 224))
77     face = img_to_array(face)
78     face = preprocess_input(face)
79     face = np.expand_dims(face, axis=0)

```

- The image's region of interest is extracted, i.e., the face. We then apply other operations on it.
- We change the order of the colour channels and resize it and then pre-process it.
- Expand_dims is a NumPy function which inserts an additional axis at the specified position. This axis stores the probability of the person wearing a mask or not.

7) Display the Result on the Screen

```

81     # pass the face through the model to determine if the face
82     # has a mask or not
83     (mask, withoutMask) = model.predict(face)[0]
84
85     # determine the class label and color we'll use to draw
86     # the bounding box and text
87     label = "Mask" if mask > withoutMask else "No Mask"
88     color = (0, 255, 0) if label == "Mask" else (0, 0, 255)
89

```

- The model predicts whether the person is wearing a face mask or not.
- We compare the probabilities and build a label which will be displayed on the screen.

```

90     # include the probability in the label
91     label = "{}: {:.2f}%".format(label, max(mask, withoutMask) * 100)
92
93     # display the label and bounding box rectangle on the output
94     # frame
95     cv2.putText(image, label, (startX, startY - 10),
96                 cv2.FONT_HERSHEY_SIMPLEX, 0.45, color, 2)
97     cv2.rectangle(image, (startX, startY), (endX, endY), color, 2)
98
99     # show the output image
100    cv2.imshow("Output", image)
101    cv2.waitKey(0)

```

- We display a rectangular box around the face with a label – mask or no mask. There is also a confidence label which displays how much the model is confident of the result out of 100.
- The interface waits for the user to press any key to exit runtime.

The complete code without explanation is available in Appendix-A.

9.3 Detection of Face Mask during a Live Video Stream

1) Importing Libraries

```

2      # python detect_mask_video.py
3
4      # import the necessary packages
5      from tensorflow.keras.applications.mobilenet_v2 import preprocess_input
6      from tensorflow.keras.preprocessing.image import img_to_array
7      from tensorflow.keras.models import load_model
8      from imutils.video import VideoStream
9      import numpy as np
10     import argparse
11     import imutils
12     import time
13     import cv2
14     import os

```

- Import the necessary packages and dependencies.
- For using the webcam we require VideoStream. This allows us to capture the frames.
- We also require the time library.

2) Processing the Video frame

```

16     def detect_and_predict_mask(frame, faceNet, maskNet):
17         # grab the dimensions of the frame and then construct a blob
18         # from it
19         (h, w) = frame.shape[:2]
20         blob = cv2.dnn.blobFromImage(frame, 1.0, (300, 300),
21                                     (104.0, 177.0, 123.0))

```

- We create a function with arguments:
 - i) Frame: The particular frame captured by our webcam.
 - ii) faceNet: Face detection.
 - iii) maskNet: Mask detection.
- We grab the spatial dimensions of the frame.
- We create a blob as discussed before.

3) Face Detections

```

23         # pass the blob through the network and obtain the face detections
24         faceNet.setInput(blob)
25         detections = faceNet.forward()
26
27         # initialize our list of faces, their corresponding locations,
28         # and the list of predictions from our face mask network
29         faces = []
30         locs = []
31         preds = []

```

- We obtain the face detection by passing it through the network.
- We initialize list of faces, their corresponding locations, and the list of predictions from face mask network.

The complete code without explanation is available in Appendix-A.

4) Building Confidence Threshold

```

33      # loop over the detections
34      for i in range(0, detections.shape[2]):
35          # extract the confidence (i.e., probability) associated with
36          # the detection
37          confidence = detections[0, 0, i, 2]
38
39          # filter out weak detections by ensuring the confidence is
40          # greater than the minimum confidence
41          if confidence > args["confidence"]:
42              # compute the (x, y)-coordinates of the bounding box for
43              # the object
44              box = detections[0, 0, i, 3:7] * np.array([w, h, w, h])
45              (startX, startY, endX, endY) = box.astype("int")

```

- As discussed before in 'Building a Confidence Threshold for Face Detection' (page 39) we let go of images with confidence lower than the threshold.
- We create a box with the specified dimensions.

5) Pre-processing the Input Frames

```

47      # ensure the bounding boxes fall within the dimensions of
48      # the frame
49      (startX, startY) = (max(0, startX), max(0, startY))
50      (endX, endY) = (min(w - 1, endX), min(h - 1, endY))
51
52      # extract the face ROI, convert it from BGR to RGB channel
53      # ordering, resize it to 224x224, and preprocess it
54      face = frame[startY:endY, startX:endX]
55      face = cv2.cvtColor(face, cv2.COLOR_BGR2RGB)
56      face = cv2.resize(face, (224, 224))
57      face = img_to_array(face)
58      face = preprocess_input(face)
59      face = np.expand_dims(face, axis=0)
60

```

- The same operations as in 'Pre-processing the Input Image' (page 40).

```

61      # add the face and bounding boxes to their respective
62      # lists
63      faces.append(face)
64      locs.append((startX, startY, endX, endY))
65

```

- Since there are multiple frames in a video, we append each frame's operation in the lists we initialized before.

6) Face mask Detection and Prediction

```

66         # only make a predictions if at least one face was detected
67         if len(faces) > 0:
68             # for faster inference we'll make batch predictions on *all*
69             # faces at the same time rather than one-by-one predictions
70             # in the above `for` loop
71             preds = maskNet.predict(faces)
72
73         # return a 2-tuple of the face locations and their corresponding
74         # locations
75         return (locs, preds)

```

- We make a prediction of face-mask detection only if there were one or more face detections. We use the maskNet network.
- We return the location of the faces and the prediction based on it.

7) Argparsing Arguments

```

77     # construct the argument parser and parse the arguments
78     ap = argparse.ArgumentParser()
79     ap.add_argument("-f", "--face", type=str,
80                    default="face_detector",
81                    help="path to face detector model directory")
82     ap.add_argument("-m", "--model", type=str,
83                    default="mask_detector.model",
84                    help="path to trained face mask detector model")
85     ap.add_argument("-c", "--confidence", type=float, default=0.5,
86                    help="minimum probability to filter weak detections")
87     args = vars(ap.parse_args())

```

- Argument parsing like discussed earlier. The arguments parsed are the same as before.

8) Load the Face and Face Mask Detection Models

```

89     # Load our serialized face detector model from disk
90     print("[INFO] loading face detector model...")
91     prototxtPath = os.path.sep.join([args["face"], "deploy.prototxt"])
92     weightsPath = os.path.sep.join([args["face"],
93                                     "res10_300x300_ssd_iter_140000.caffemodel"])
94     faceNet = cv2.dnn.readNet(prototxtPath, weightsPath)
95
96     # Load the face mask detector model from disk
97     print("[INFO] loading face mask detector model...")
98     maskNet = load_model(args["model"])
99

```

- We load the face detector and face-mask detector models.

9) Start the Video Stream through WebCam

```

100 # initialize the video stream and allow the camera sensor to warm up
101 print("[INFO] starting video stream...")
102 vs = VideoStream(src=0).start()
103 time.sleep(2.0)

```

- We start our video stream through our webcam with delay of 2 seconds.

```

105 # Loop over the frames from the video stream
106 while True:
107     # grab the frame from the threaded video stream and resize it
108     # to have a maximum width of 400 pixels
109     frame = vs.read()
110     frame = imutils.resize(frame, width=400)
111
112     # detect faces in the frame and determine if they are wearing a
113     # face mask or not
114     (locs, preds) = detect_and_predict_mask(frame, faceNet, maskNet)

```

- We grab the frame from the threaded video stream and resize it to have a maximum width of 400 pixels.
- We call the function `detect_and_predict_mask` that we defined earlier and store the return values in `locs` and `preds`.

10) Prediction for Face Mask

```

116 # Loop over the detected face locations and their corresponding
117 # locations
118 for (box, pred) in zip(locs, preds):
119     # unpack the bounding box and predictions
120     (startX, startY, endX, endY) = box
121     (mask, withoutMask) = pred
122
123     # determine the class label and color we'll use to draw
124     # the bounding box and text
125     label = "Mask" if mask > withoutMask else "No Mask"
126     color = (0, 255, 0) if label == "Mask" else (0, 0, 255)
127
128     # include the probability in the label
129     label = "{}: {:.2f}%".format(label, max(mask, withoutMask) * 100)
130

```

- We do the same operations as we did `face_mask_detection` with the only difference here being that we are operating on a single frame of a video rather than an image.
- We define a label as we did before.

```

131     # display the label and bounding box rectangle on the output
132     # frame
133     cv2.putText(frame, label, (startX, startY - 10),
134                 cv2.FONT_HERSHEY_SIMPLEX, 0.45, color, 2)
135     cv2.rectangle(frame, (startX, startY), (endX, endY), color, 2)
136
137     # show the output frame
138     cv2.imshow("Frame", frame)
139     key = cv2.waitKey(1) & 0xFF
140
141     # if the `q` key was pressed, break from the loop
142     if key == ord("q"):
143         break
144
145     # do a bit of cleanup
146     cv2.destroyAllWindows()
147     vs.stop()

```

- Like before the label and the prediction accuracy is displayed alongside the rectangle.
- The output frame is shown in real time. As soon as the 'q' key is pressed the video stream stops and we exit runtime.

CHAPTER 10

Loopholes in the Model

- Our model is not able to detect faces which are largely covered with masks. The mask obscures with the face-detection process and hence it is not able to pass the confidence threshold.
- The model is only able to detect light-coloured masks and fails to detect other dark colours. This is due to the training dataset which only contains light-coloured masks.
- The model is susceptible to over-fitting. This happens when the model is generalized to the training dataset and fails to classify the testing dataset correctly. This happens when the dataset is not large.

How can the Model be Improved

- Insert more layers into the neural network. More layers will help to recognize parts of the face better and hence the model would be able to have high confidence value on the face detections.
- Improve the dataset by increasing the diversity of the training images. Artificial datasets could hamper the performance of the neural network.
- Increase the number of training data and apply fine-tuning. DropOuts, Regularization are some of the techniques through which we could fit our data better in the model.

Conclusion

Using the above techniques and functions we can train a deep neural net that can detect whether a person is wearing a face mask or not.

APPENDIX-A

train_mask_detector.py

```

4  # import the necessary packages
5  from tensorflow.keras.preprocessing.image import ImageDataGenerator
6  from tensorflow.keras.applications import MobileNetV2
7  from tensorflow.keras.layers import AveragePooling2D
8  from tensorflow.keras.layers import Dropout
9  from tensorflow.keras.layers import Flatten
10 from tensorflow.keras.layers import Dense
11 from tensorflow.keras.layers import Input
12 from tensorflow.keras.models import Model
13 from tensorflow.keras.optimizers import Adam
14 from tensorflow.keras.applications.mobilenet_v2 import preprocess_input
15 from tensorflow.keras.preprocessing.image import img_to_array
16 from tensorflow.keras.preprocessing.image import load_img
17 from tensorflow.keras.utils import to_categorical
18 from sklearn.preprocessing import LabelBinarizer
19 from sklearn.model_selection import train_test_split
20 from sklearn.metrics import classification_report
21 from imutils import paths
22 import matplotlib.pyplot as plt
23 import numpy as np
24 import argparse
25 import os
26
27 # construct the argument parser and parse the arguments
28 ap = argparse.ArgumentParser()
29 ap.add_argument("-d", "--dataset", required=True,
30                 help="path to input dataset")
31 ap.add_argument("-p", "--plot", type=str, default="plot.png",
32                 help="path to output loss/accuracy plot")
33 ap.add_argument("-m", "--model", type=str,
34                 default="mask_detector.model",
35                 help="path to output face mask detector model")
36 args = vars(ap.parse_args())

```

The complete code without explanation is available in Appendix-A.

```

38 # initialize the initial learning rate, number of epochs to train for,
39 # and batch size
40 INIT_LR = 1e-4
41 EPOCHS = 20
42 BS = 32
43
44 # grab the list of images in our dataset directory, then initialize
45 # the list of data (i.e., images) and class images
46 print("[INFO] loading images...")
47 imagePath = list(paths.list_images(args["dataset"]))
48 data = []
49 labels = []
50
51 # loop over the image paths
52 for imagePath in imagePath:
53     # extract the class label from the filename
54     label = imagePath.split(os.path.sep)[-2]
55
56     # load the input image (224x224) and preprocess it
57     image = load_img(imagePath, target_size=(224, 224))
58     image = img_to_array(image)
59     image = preprocess_input(image)
60
61     # update the data and labels lists, respectively
62     data.append(image)
63     labels.append(label)
64
65 # convert the data and labels to NumPy arrays
66 data = np.array(data, dtype="float32")
67 labels = np.array(labels)
68
69 # perform one-hot encoding on the labels
70 lb = LabelBinarizer()
71 labels = lb.fit_transform(labels)
72 labels = to_categorical(labels)
73

```

```

74 # partition the data into training and testing splits using 75% of
75 # the data for training and the remaining 25% for testing
76 (trainX, testX, trainY, testY) = train_test_split(data, labels,
77     test_size=0.20, stratify=labels, random_state=42)
78
79 # construct the training image generator for data augmentation
80 aug = ImageDataGenerator(
81     rotation_range=20,
82     zoom_range=0.15,
83     width_shift_range=0.2,
84     height_shift_range=0.2,
85     shear_range=0.15,
86     horizontal_flip=True,
87     fill_mode="nearest")
88
89 # Load the MobileNetV2 network, ensuring the head FC layer sets are
90 # left off
91 baseModel = MobileNetV2(weights="imagenet", include_top=False,
92     input_tensor=Input(shape=(224, 224, 3)))
93
94 # construct the head of the model that will be placed on top of the
95 # the base model
96 headModel = baseModel.output
97 headModel = AveragePooling2D(pool_size=(7, 7))(headModel)
98 headModel = Flatten(name="flatten")(headModel)
99 headModel = Dense(128, activation="relu")(headModel)
100 headModel = Dropout(0.5)(headModel)
101 headModel = Dense(2, activation="softmax")(headModel)
102
103 # place the head FC model on top of the base model (this will become
104 # the actual model we will train)
105 model = Model(inputs=baseModel.input, outputs=headModel)
106
107 # loop over all layers in the base model and freeze them so they will
108 # *not* be updated during the first training process
109 for layer in baseModel.layers:
110     layer.trainable = False

```



```

112 # compile our model
113 print("[INFO] compiling model...")
114 opt = Adam(lr=INIT_LR, decay=INIT_LR / EPOCHS)
115 model.compile(loss="binary_crossentropy", optimizer=opt,
116               metrics=["accuracy"])
117
118 # train the head of the network
119 print("[INFO] training head...")
120 H = model.fit(
121     aug.flow(trainX, trainY, batch_size=BS),
122     steps_per_epoch=len(trainX) // BS,
123     validation_data=(testX, testY),
124     validation_steps=len(testX) // BS,
125     epochs=EPOCHS)
126
127 # make predictions on the testing set
128 print("[INFO] evaluating network...")
129 predIdxs = model.predict(testX, batch_size=BS)
130
131 # for each image in the testing set we need to find the index of the
132 # label with corresponding largest predicted probability
133 predIdxs = np.argmax(predIdxs, axis=1)
134
135 # show a nicely formatted classification report
136 print(classification_report(testY.argmax(axis=1), predIdxs,
137                             target_names=lb.classes_))
138
139 # serialize the model to disk
140 print("[INFO] saving mask detector model...")
141 model.save(args["model"], save_format="h5")
142
143 # plot the training loss and accuracy
144 N = EPOCHS
145 plt.style.use("ggplot")
146 plt.figure()
147 plt.plot(np.arange(0, N), H.history["loss"], label="train_loss")
148 plt.plot(np.arange(0, N), H.history["val_loss"], label="val_loss")
149 plt.plot(np.arange(0, N), H.history["accuracy"], label="train_acc")
150 plt.plot(np.arange(0, N), H.history["val_accuracy"], label="val_acc")
151 plt.title("Training Loss and Accuracy")
152 plt.xlabel("Epoch #")
153 plt.ylabel("Loss/Accuracy")
154 plt.legend(loc="lower left")
155 plt.savefig(args["plot"])

```


`detect_mask_image.py`

```

4  # import the necessary packages
5  from tensorflow.keras.applications.mobilenet_v2 import preprocess_input
6  from tensorflow.keras.preprocessing.image import img_to_array
7  from tensorflow.keras.models import load_model
8  import numpy as np
9  import argparse
10 import cv2
11 import os
12
13 # construct the argument parser and parse the arguments
14 ap = argparse.ArgumentParser()
15 ap.add_argument("-i", "--image", required=True,
16                 help="path to input image")
17 ap.add_argument("-f", "--face", type=str,
18                 default="face_detector",
19                 help="path to face detector model directory")
20 ap.add_argument("-m", "--model", type=str,
21                 default="mask_detector.model",
22                 help="path to trained face mask detector model")
23 ap.add_argument("-c", "--confidence", type=float, default=0.5,
24                 help="minimum probability to filter weak detections")
25 args = vars(ap.parse_args())
26
27 # Load our serialized face detector model from disk
28 print("[INFO] loading face detector model...")
29 prototxtPath = os.path.sep.join([args["face"], "deploy.prototxt"])
30 weightsPath = os.path.sep.join([args["face"],
31                                 "res10_300x300_ssd_iter_140000.caffemodel"])
32 net = cv2.dnn.readNet(prototxtPath, weightsPath)
33
34 # Load the face mask detector model from disk
35 print("[INFO] loading face mask detector model...")
36 model = load_model(args["model"])
37
38 # Load the input image from disk, clone it, and grab the image spatial
39 # dimensions
40 image = cv2.imread(args["image"])
41 orig = image.copy()
42 (h, w) = image.shape[:2]
43

```

```

44     # construct a blob from the image
45     blob = cv2.dnn.blobFromImage(image, 1.0, (300, 300),
46                                  (104.0, 177.0, 123.0))
47
48     # pass the blob through the network and obtain the face detections
49     print("[INFO] computing face detections...")
50     net.setInput(blob)
51     detections = net.forward()
52
53     # loop over the detections
54     for i in range(0, detections.shape[2]):
55         # extract the confidence (i.e., probability) associated with
56         # the detection
57         confidence = detections[0, 0, i, 2]
58
59         # filter out weak detections by ensuring the confidence is
60         # greater than the minimum confidence
61         if confidence > args["confidence"]:
62             # compute the (x, y)-coordinates of the bounding box for
63             # the object
64             box = detections[0, 0, i, 3:7] * np.array([w, h, w, h])
65             (startX, startY, endX, endY) = box.astype("int")
66
67             # ensure the bounding boxes fall within the dimensions of
68             # the frame
69             (startX, startY) = (max(0, startX), max(0, startY))
70             (endX, endY) = (min(w - 1, endX), min(h - 1, endY))
71
72             # extract the face ROI, convert it from BGR to RGB channel
73             # ordering, resize it to 224x224, and preprocess it
74             face = image[startY:endY, startX:endX]
75             face = cv2.cvtColor(face, cv2.COLOR_BGR2RGB)
76             face = cv2.resize(face, (224, 224))
77             face = img_to_array(face)
78             face = preprocess_input(face)
79             face = np.expand_dims(face, axis=0)

```

```
81     # pass the face through the model to determine if the face
82     # has a mask or not
83     (mask, withoutMask) = model.predict(face)[0]
84
85     # determine the class label and color we'll use to draw
86     # the bounding box and text
87     label = "Mask" if mask > withoutMask else "No Mask"
88     color = (0, 255, 0) if label == "Mask" else (0, 0, 255)
89
90     # include the probability in the label
91     label = "{}: {:.2f}%".format(label, max(mask, withoutMask) * 100)
92
93     # display the label and bounding box rectangle on the output
94     # frame
95     cv2.putText(image, label, (startX, startY - 10),
96                 cv2.FONT_HERSHEY_SIMPLEX, 0.45, color, 2)
97     cv2.rectangle(image, (startX, startY), (endX, endY), color, 2)
98
99     # show the output image
100    cv2.imshow("Output", image)
101    cv2.waitKey(0)
```

detect_mask_video.py

```
4  # import the necessary packages
5  from tensorflow.keras.applications.mobilenet_v2 import preprocess_input
6  from tensorflow.keras.preprocessing.image import img_to_array
7  from tensorflow.keras.models import load_model
8  from imutils.video import VideoStream
9  import numpy as np
10 import argparse
11 import imutils
12 import time
13 import cv2
14 import os
15
16 def detect_and_predict_mask(frame, faceNet, maskNet):
17     # grab the dimensions of the frame and then construct a blob
18     # from it
19     (h, w) = frame.shape[:2]
20     blob = cv2.dnn.blobFromImage(frame, 1.0, (300, 300),
21     |     (104.0, 177.0, 123.0))
22
23     # pass the blob through the network and obtain the face detections
24     faceNet.setInput(blob)
25     detections = faceNet.forward()
26
27     # initialize our list of faces, their corresponding locations,
28     # and the list of predictions from our face mask network
29     faces = []
30     locs = []
31     preds = []
```

```

33     # Loop over the detections
34     for i in range(0, detections.shape[2]):
35         # extract the confidence (i.e., probability) associated with
36         # the detection
37         confidence = detections[0, 0, i, 2]
38
39         # filter out weak detections by ensuring the confidence is
40         # greater than the minimum confidence
41         if confidence > args["confidence"]:
42             # compute the (x, y)-coordinates of the bounding box for
43             # the object
44             box = detections[0, 0, i, 3:7] * np.array([w, h, w, h])
45             (startX, startY, endX, endY) = box.astype("int")
46
47             # ensure the bounding boxes fall within the dimensions of
48             # the frame
49             (startX, startY) = (max(0, startX), max(0, startY))
50             (endX, endY) = (min(w - 1, endX), min(h - 1, endY))
51
52             # extract the face ROI, convert it from BGR to RGB channel
53             # ordering, resize it to 224x224, and preprocess it
54             face = frame[startY:endY, startX:endX]
55             face = cv2.cvtColor(face, cv2.COLOR_BGR2RGB)
56             face = cv2.resize(face, (224, 224))
57             face = img_to_array(face)
58             face = preprocess_input(face)
59             face = np.expand_dims(face, axis=0)
60
61             # add the face and bounding boxes to their respective
62             # lists
63             faces.append(face)
64             locs.append((startX, startY, endX, endY))
65

```

```

66         # only make a predictions if at least one face was detected
67         if len(faces) > 0:
68             # for faster inference we'll make batch predictions on *all*
69             # faces at the same time rather than one-by-one predictions
70             # in the above `for` loop
71             preds = maskNet.predict(faces)
72
73         # return a 2-tuple of the face locations and their corresponding
74         # locations
75         return (locs, preds)
76
77     # construct the argument parser and parse the arguments
78     ap = argparse.ArgumentParser()
79     ap.add_argument("-f", "--face", type=str,
80                     default="face_detector",
81                     help="path to face detector model directory")
82     ap.add_argument("-m", "--model", type=str,
83                     default="mask_detector.model",
84                     help="path to trained face mask detector model")
85     ap.add_argument("-c", "--confidence", type=float, default=0.5,
86                     help="minimum probability to filter weak detections")
87     args = vars(ap.parse_args())
88
89     # load our serialized face detector model from disk
90     print("[INFO] loading face detector model...")
91     prototxtPath = os.path.sep.join([args["face"], "deploy.prototxt"])
92     weightsPath = os.path.sep.join([args["face"],
93                                     "res10_300x300_ssd_iter_140000.caffemodel"])
94     faceNet = cv2.dnn.readNet(prototxtPath, weightsPath)
95
96     # load the face mask detector model from disk
97     print("[INFO] loading face mask detector model...")
98     maskNet = load_model(args["model"])
99
100    # initialize the video stream and allow the camera sensor to warm up
101    print("[INFO] starting video stream...")
102    vs = VideoStream(src=0).start()
103    time.sleep(2.0)

```

```

105     # loop over the frames from the video stream
106     while True:
107         # grab the frame from the threaded video stream and resize it
108         # to have a maximum width of 400 pixels
109         frame = vs.read()
110         frame = imutils.resize(frame, width=400)
111
112         # detect faces in the frame and determine if they are wearing a
113         # face mask or not
114         (locs, preds) = detect_and_predict_mask(frame, faceNet, maskNet)
115
116         # loop over the detected face locations and their corresponding
117         # locations
118         for (box, pred) in zip(locs, preds):
119             # unpack the bounding box and predictions
120             (startX, startY, endX, endY) = box
121             (mask, withoutMask) = pred
122
123             # determine the class label and color we'll use to draw
124             # the bounding box and text
125             label = "Mask" if mask > withoutMask else "No Mask"
126             color = (0, 255, 0) if label == "Mask" else (0, 0, 255)
127
128             # include the probability in the label
129             label = "{}: {:.2f}%".format(label, max(mask, withoutMask) * 100)
130
131             # display the label and bounding box rectangle on the output
132             # frame
133             cv2.putText(frame, label, (startX, startY - 10),
134                 cv2.FONT_HERSHEY_SIMPLEX, 0.45, color, 2)
135             cv2.rectangle(frame, (startX, startY), (endX, endY), color, 2)
136
137         # show the output frame
138         cv2.imshow("Frame", frame)
139         key = cv2.waitKey(1) & 0xFF
140
141         # if the `q` key was pressed, break from the loop
142         if key == ord("q"):
143             break
144
145     # do a bit of cleanup
146     cv2.destroyAllWindows()
147     vs.stop()

```


APPENDIX-B

Argparse and Command Line Arguments in Python

Command line arguments are flags given to a program/script at runtime. They contain additional information for our program so that it can execute. This makes the program more user interactive and hence it is used very often. This allows us to give our program different input on the fly without changing the code.

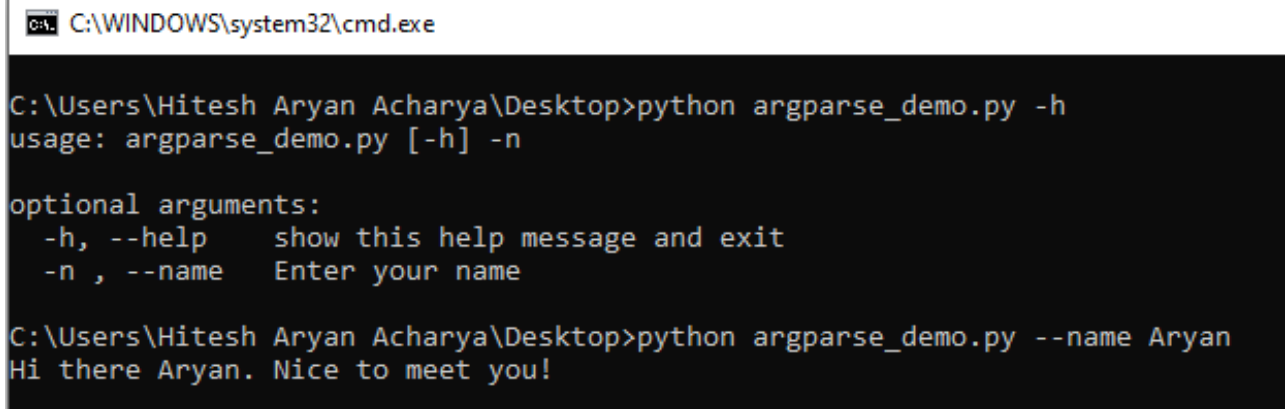
Command line arguments can be inculcated using argparsing. Argparsing is used a lot in computer vision and deep learning. We'll be using argparse in the following and hence it is imperative that we understand about it.

```
1
2 #demo to illustarte function of argparse
3 import argparse
4
5 ap = argparse.ArgumentParser()
6 ap.add_argument('-n', '--name', type = str, metavar = '', required = True, help = 'Enter your name')
7 args = vars(ap.parse_args())
8
9 print("Hi there {}. Nice to meet you!".format(args["name"]))
10
```

Figure 45- Example illustrating Argument Parsing

- We first import the library argparse.
- We add argument using the in-built function add_argument. '-n', '--name' is the short and long notation for the argument.
- We define the parameters. The type here is 'string' which means the argument to be parsed is of the type string.
- 'required = True' means the argument is necessary. If the user runs the program without the argument it'll show an error.
- The help option can be used by the user to see the structure of argument parsing. Below we use the '-h' command to call help.

- We create a dictionary 'args' which stores all the arguments parsed by the user. It is helpful when there are multiple arguments.



```
C:\WINDOWS\system32\cmd.exe

C:\Users\Hitesh Aryan Acharya\Desktop>python argparse_demo.py -h
usage: argparse_demo.py [-h] -n

optional arguments:
  -h, --help      show this help message and exit
  -n, --name      Enter your name

C:\Users\Hitesh Aryan Acharya\Desktop>python argparse_demo.py --name Aryan
Hi there Aryan. Nice to meet you!
```

Figure 46- Parsing my name as an argument

APPENDIX-C

Pre-processing the image Input and MobileNetV2 architecture

Whenever working with neural networks it is imperative that we pre-process the input. Data pre-processing aims to facilitate the training/testing process by appropriately transforming and scaling the entire dataset. Pre-processing is necessary before training the machine learning models. Pre-processing removes outliers and scales the features to an equivalent range. In our deep learning model as well, we will be pre-processing our dataset. Pre-processing increases efficiency and reduces time consumption.

blobFromImage

In the context of deep learning and image classification, these pre-processing tasks normally involve:

1. Mean subtraction
2. Scaling
3. Optionally channel swapping

All these processes can be implemented using the function `cv2.dnn.blobFromImage`. This is an in-built function in the OpenCV library and is used extensively during object detection. We will go through each technique.

Mean Subtraction and Scaling: Mean subtraction is used to help combat illumination changes in the input images in our dataset. We can therefore view mean subtraction as a technique used to aid our Convolutional Neural Networks.

We basically compute the mean across all the colour channels in our input image. The mean is stored in a python tuple (three entries, one each for red, blue and green). The standard deviation or the scaling factor is also calculated across the entire image. After this we perform the following operation:

1. $R = (R - \mu) / \sigma$
2. $B = (B - \mu) / \sigma$
3. $G = (G - \mu) / \sigma$

We have hence normalized the image. We do all this using a simple function,

```
blob = cv2.dnn.blobFromImage(image, scalefactor=1.0, size, mean, swapRB=True)
```

--**image**: This is the image we want to preprocess.

--**scalefactor**: σ in the above equations is the scalefactor.

--**size**: Here we supply the spatial size that the Convolutional Neural Network expects.

--**mean**: This is the tuple containing the means of the three channels.

--**swapRB**: OpenCV assumes images are in BGR channel order; however, the `mean` value assumes we are using RGB order. To resolve this discrepancy we can swap the R and B channels in image by setting this value to `True`. By default OpenCV performs this channel swapping for us.

The `cv2.dnn.blobFromImage` function returns a blob which is our input image after mean subtraction, normalizing, and channel swapping.

MobileNetV2 architecture

In 2017 Google introduced MobileNetV1, a family of general purpose computer vision neural networks designed with mobile devices in mind to support classification, detection and more. The ability to run deep networks on personal mobile devices improves user experience, offering anytime, anywhere access, with additional benefits for security, privacy, and energy consumption. MobileNetV2 is a significant improvement over MobileNetV1 and pushes the state of the art for mobile visual recognition including classification, object detection and semantic segmentation.

So in a nutshell MobilenetV2 is a highly efficient architecture that can be applied to embedded devices with limited computational efficiency (e.g. Raspberry Pi, NVIDIA Jetson Nano etc.).

Deploying our face mask detector to embedded devices could reduce the cost of manufacturing such face mask detection systems, hence why we choose to use this architecture.

APPENDIX-D

LabelBinarizer, fit_transform and to_categorical

1. These concepts are best described using examples. LabelBinarizer is a technique to classify the labelled data as seen from the example below:

```
from numpy import array
from sklearn.preprocessing import LabelBinarizer

# define example
data = ['cold', 'cold', 'warm', 'cold', 'hot', 'hot', 'warm', 'cold',
        'warm', 'hot']
values = array(data)
print "Data: ", values
#Binary encode
lb = LabelBinarizer()
print "Label Binarizer:", lb.fit_transform(values)
```

Figure 47- Example illustrating LabelBinarizer

```
Data: ['cold' 'cold' 'warm' 'cold' 'hot' 'hot' 'warm' 'cold' 'warm' 'hot']
Label Binarizer: [[1 0 0]
 [1 0 0]
 [0 1 0]
 [1 0 0]
 [0 1 0]
 [0 1 0]
 [0 0 1]
 [1 0 0]
 [0 0 1]
 [0 1 0]]
```

Figure 48- The format of the labelled output has changed

2. to_categorical function converts a class vector(integers) to binary class matrix.

```
>>> a = tf.keras.utils.to_categorical([0, 1, 2, 3], num_classes=4)
>>> a = tf.constant(a, shape=[4, 4])
>>> print(a)
tf.Tensor(
  [[1. 0. 0. 0.]
   [0. 1. 0. 0.]
   [0. 0. 1. 0.]
   [0. 0. 0. 1.]], shape=(4, 4), dtype=float32)
```

Figure 49- Using the to_categorical function

3. To center the data (make it have zero mean and unit standard error), we subtract the mean and then divide the result by the standard deviation:

$$x' = \frac{(x - \mu)}{\sigma}$$

We do that on the training set of data. But then we have to apply the same transformation to our testing set (e.g. in cross-validation), or to newly obtained examples before forecast. But we have to use the exact same two parameters μ and σ (values) that we used for centering the training set.

Hence, sklearn's (library) transform's `fit()` just calculates the parameters and saves them as an internal object's state. Afterwards, we can call its `transform()` method to apply the transformation to any particular set of examples.

`fit_transform()` joins these two steps and is used for the initial fitting of parameters on the training set x , while also returning the transformed x' . Internally, the transformer object just calls first `fit()` and then `transform()` on the same data.

APPENDIX-E

How to Run the Face-Mask Detector Model on Windows Terminal

Install Python

- On the terminal pass the command `python --version`. Any version above python 3.5 is good enough to run the program.
- Next pass the command `pip --version`. The version of pip alongside the python version will be displayed if pip was installed as a package. [Pip 20.0.2](#) is the latest version of pip (as of 13 Jun. 20).
- If you have an older version of python installed, it is recommended that you uninstall the python package and proceed to download the latest version from the net.
- To uninstall, go to Control Panel→Programs→Programs and Features→Uninstall a program. Choose the python package you want to uninstall.
- Go to <https://www.python.org/downloads/> and click on Download Python 3.8.3 (as of 13 June 2020).
- Follow the instructions and click on install.
- The latest version of python alongside the latest pip will be installed on your desktop.
- Now go to terminal and repeat the first step. You'll see the latest version of python and pip installed.

```
Microsoft Windows [Version 10.0.18362.900]
(c) 2019 Microsoft Corporation. All rights reserved.

C:\Users\Hitesh Aryan Acharya>python --version
Python 3.7.3

C:\Users\Hitesh Aryan Acharya>pip --version
pip 20.0.2 from C:\Users\Hitesh Aryan Acharya\Anaconda3\lib\site-packages\pip (python 3.7)
```

Figure 50- Installing the latest python version

Importing the necessary Libraries and Packages

- The code involves TensorFlow and keras and other important math packages. We first need to use pip to install these packages.
- To install TensorFlow, pass the command:
`pip install --upgrade tensorflow`
- Now we need to install some python dependencies. Pass the following commands:
`pip install numpy scipy`
`pip install scikit-learn`
`pip install pillow`
`pip install h5py`
`pip install imutils`
- Now to install keras, pass the following command:
`pip install keras`

Extracting the Files

- Extract all the files from the zip folder. The extracted folder will be named as face_mask_detector.
- Copy the extracted folder to desktop.
- Go to terminal and pass the command `cd desktop/face_mask_detector`.

Name	Date modified	Type	Size
dataset	12-06-2020 17:06	File folder	
examples	23-05-2020 19:05	File folder	
face_detector	23-05-2020 19:05	File folder	
detect_mask_image	23-05-2020 19:05	JetBrains PyCharm	4 KB
detect_mask_video	23-05-2020 19:05	JetBrains PyCharm	5 KB
model	12-06-2020 18:06	File	11,215 KB
plot	12-06-2020 18:07	PNG File	24 KB
train_mask_detector	06-06-2020 00:06	JetBrains PyCharm	6 KB

Figure 51- face_mask_detector directory

- The directory contains the above files.
- The dataset folder contains the training and testing examples (images) for the neural network.

- Examples folder contains the live examples we'll be using to test the neural network.
- The face_detector folder contains the face classifier which will be used to detect faces in images.
- detect_mask_image, detect_mask_video and train_mask_detector are python files which we will be running on our terminal.
- model is a text file describing the model of the deep learning structure.
- plot is a image file depicting the training and validation loss and accuracy of the model (during train_mask_detector).

Running the Face-Mask Detector

- On the terminal pass the command:

```
python train_mask_detector.py --dataset dataset
```

The above command runs the python code in train_mask_detector using the dataset folder.

- After the training is completed, we can test our network using the examples in the examples folder. To do so pass the command:

```
python detect_mask_image.py --image
examples/example_01.png
```

This will display a picture with a man wearing a face-mask. As we can see our model is able to identify the face mask. The classifier also labels the image correctly.

- Similarly we can check for other examples. Pass the command:

```
python detect_mask_image.py --image
examples/example_02.png python
detect_mask_image.py --image
examples/example_03.png
```

- We can also use the existing model for face-mask detection during a live video stream. To do this pass the command:

```
python detect_mask_video.py
```

Your webcam will now start up. With good accuracy the model will predict whether you're wearing a face mask or not. Good lighting is required for optimal performance.

APPENDIX-F

Some snips during runtime

```

34/34 [=====] - 47s 1s/step - loss: 0.4336 - accuracy: 0.8098 - val_loss: 0.1668 - val_accuracy: 0.9689
Epoch 3/20
34/34 [=====] - 47s 1s/step - loss: 0.3141 - accuracy: 0.8755 - val_loss: 0.1255 - val_accuracy: 0.9688
Epoch 4/20
34/34 [=====] - 45s 1s/step - loss: 0.2616 - accuracy: 0.8951 - val_loss: 0.1006 - val_accuracy: 0.9727
Epoch 5/20
34/34 [=====] - 48s 1s/step - loss: 0.2382 - accuracy: 0.9064 - val_loss: 0.1163 - val_accuracy: 0.9648
Epoch 6/20
34/34 [=====] - 51s 1s/step - loss: 0.2174 - accuracy: 0.9127 - val_loss: 0.0798 - val_accuracy: 0.9805
Epoch 7/20
34/34 [=====] - 50s 1s/step - loss: 0.1861 - accuracy: 0.9280 - val_loss: 0.0736 - val_accuracy: 0.9805
Epoch 8/20
34/34 [=====] - 59s 2s/step - loss: 0.1787 - accuracy: 0.9298 - val_loss: 0.0596 - val_accuracy: 0.9961
Epoch 9/20
34/34 [=====] - 54s 2s/step - loss: 0.1701 - accuracy: 0.9354 - val_loss: 0.0600 - val_accuracy: 0.9885
Epoch 10/20
34/34 [=====] - 52s 2s/step - loss: 0.1544 - accuracy: 0.9430 - val_loss: 0.0480 - val_accuracy: 0.9922
Epoch 11/20
34/34 [=====] - 46s 1s/step - loss: 0.1482 - accuracy: 0.9393 - val_loss: 0.0705 - val_accuracy: 0.9805
Epoch 12/20
34/34 [=====] - 53s 2s/step - loss: 0.1527 - accuracy: 0.9363 - val_loss: 0.0514 - val_accuracy: 0.9883
Epoch 13/20
34/34 [=====] - 46s 1s/step - loss: 0.1300 - accuracy: 0.9560 - val_loss: 0.0478 - val_accuracy: 0.9922
Epoch 14/20
34/34 [=====] - 47s 1s/step - loss: 0.1251 - accuracy: 0.9478 - val_loss: 0.0390 - val_accuracy: 0.9922
Epoch 15/20
34/34 [=====] - 47s 1s/step - loss: 0.1192 - accuracy: 0.9532 - val_loss: 0.0755 - val_accuracy: 0.9766
Epoch 16/20
34/34 [=====] - 44s 1s/step - loss: 0.1133 - accuracy: 0.9569 - val_loss: 0.0424 - val_accuracy: 0.9961
Epoch 17/20
34/34 [=====] - 49s 1s/step - loss: 0.1033 - accuracy: 0.9509 - val_loss: 0.0533 - val_accuracy: 0.9844
Epoch 18/20
34/34 [=====] - 52s 2s/step - loss: 0.0990 - accuracy: 0.9607 - val_loss: 0.0358 - val_accuracy: 0.9883
Epoch 19/20
34/34 [=====] - 48s 1s/step - loss: 0.1106 - accuracy: 0.9570 - val_loss: 0.0459 - val_accuracy: 0.9883
Epoch 20/20
34/34 [=====] - 51s 1s/step - loss: 0.1006 - accuracy: 0.9607 - val_loss: 0.0487 - val_accuracy: 0.9883
[INFO] evaluating network...
      precision    recall  f1-score   support

 with_mask         0.98         1.00         0.99         138
 without_mask       1.00         0.98         0.99         138

 accuracy         0.99
 macro avg        0.99         0.99         0.99         276
 weighted avg     0.99         0.99         0.99         276


```

Figure 52- Training our neural network over 20 EPOCHS

```

=====] - 49s 1s/step - loss: 0.1033 - accuracy: 0.9509 - val_loss: 0.0533 - val_accuracy: 0.9844
=====] - 52s 2s/step - loss: 0.0990 - accuracy: 0.9607 - val_loss: 0.0358 - val_accuracy: 0.9883
=====] - 48s 1s/step - loss: 0.1106 - accuracy: 0.9570 - val_loss: 0.0459 - val_accuracy: 0.9883
[INFO] evaluating network...
precision    recall  f1-score   support
 with_mask         0.98         1.00         0.99         138
 without_mask       1.00         0.98         0.99         138
 accuracy         0.99
 macro avg        0.99         0.99         0.99         276
 weighted avg     0.99         0.99         0.99         276

```



```

resh Aryan Acharya\
ing face detector m
ing face mask detec
13:45:19.671441: I
03:45:29.678869: I
root recent call la
ect mask_image.py",
image.copy()
rort: 'NoneType' ob)
resh Aryan Acharya\
ing face detector m
ing face mask detec
23:48:10.119292: I
23:48:10.126254: I
uting face detectio
uting face detections...
resh Aryan Acharya\Desktop12-2\face_mask_detector>python detect_mask_image.py --image examples/example_01.png
ing face detector model...
ing face mask detector model...
23:50:28.832782: I tensorflow/core/platform/cpu_feature_guard.cc:142] Your CPU supports instructions that this TensorFlow binary was not compiled to use: AVX AVX2
23:50:28.839688: I tensorflow/core/common_runtime/process_util.cc:147] Creating new thread pool with default inter op setting: 2. Tune using inter_op_parallelism
uting face detections...

```

Figure 53- Detecting face-mask in an image



Figure 54- Here the model is not able to detect the face mask but correctly detects the non-masked faces

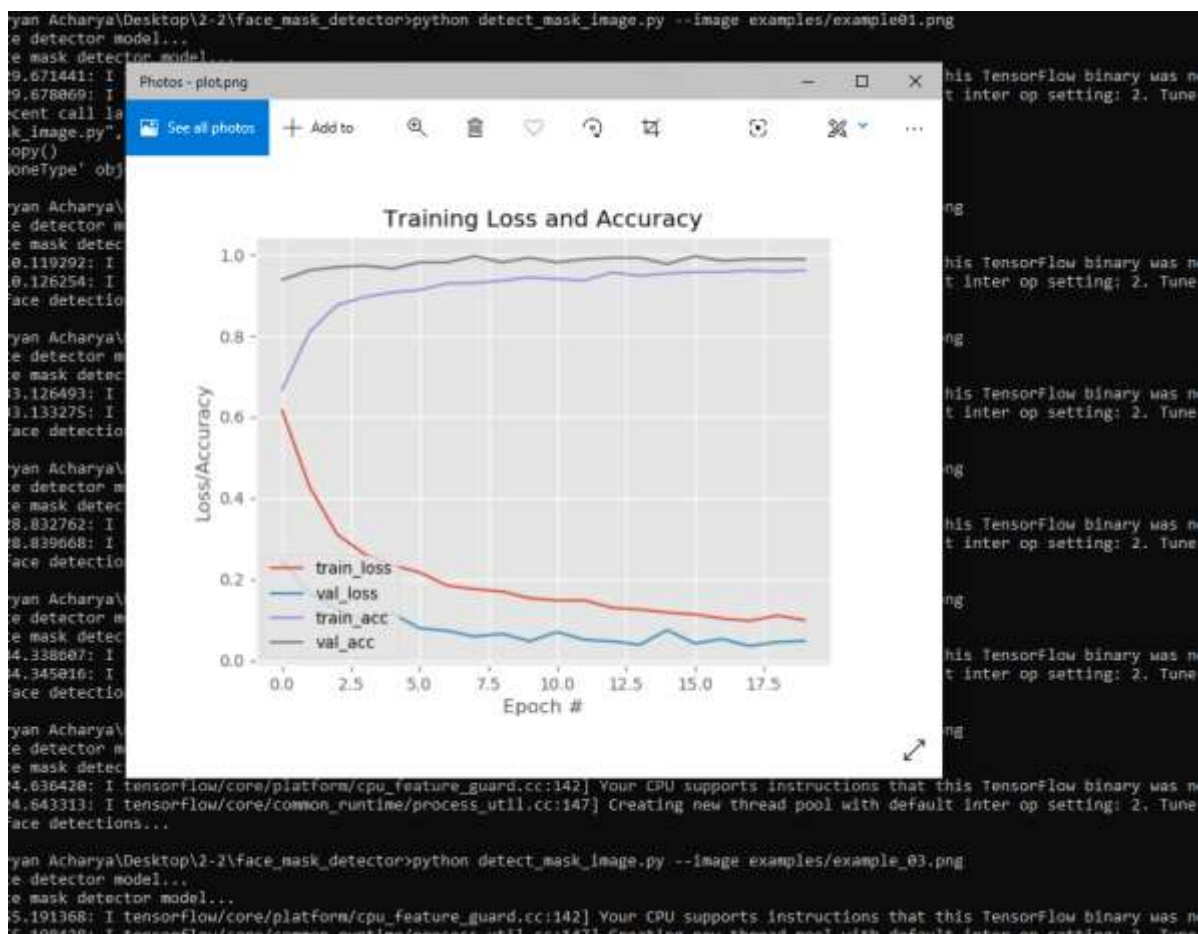


Figure 55- The plot of training and validation parameters




Figure 56- The network correctly detects my green mask during video stream




Figure 57- But it doesn't detect my dark blue mask

References

- COVID-19: Face Mask Detector with OpenCV, Keras/TensorFlow, and Deep Learning (<https://www.pyimagesearch.com/2020/05/04/covid-19-face-mask-detector-with-opencv-keras-tensorflow-and-deep-learning/>)
- Deep learning: How OpenCV's blobFromImage works (<https://www.pyimagesearch.com/2017/11/06/deep-learning-opencvs-blobfromimage-works/>)
- TensorFlow Documentation (https://www.tensorflow.org/api_docs/python/tf/keras/Model)
- Keras Documentation (<https://keras.io/documentation/>)
- Coursera- Neural Networks and Deep Learning (<https://www.coursera.org/learn/neural-networks-deep-learning/home/welcome>)
- Coursera- Introduction to TensorFlow for Artificial Intelligence, Machine Learning, and Deep Learning (<https://www.coursera.org/learn/introduction-tensorflow/home/welcome>)
- StackOverflow- Scikit-learn's LabelBinarizer (<https://stackoverflow.com/questions/50473381/scikit-learns-labelbinarizer-vs-onehotencoder>)
- StackExchange Data Science- fit_transform (<https://stackoverflow.com/questions/50473381/scikit-learns-labelbinarizer-vs-onehotencoder>)
- Confidence Threshold (<https://blog.zenggyu.com/en/post/2018-12-16/an-introduction-to-evaluation-metrics-for-object-detection>)
- MobileNetV2 architecture (<https://analyticsindiamag.com/why-googles-mobilenetv2-is-a-revolutionary-next-gen-on-device-computer-vision-network/>)
- Data Augmentation (<https://www.pyimagesearch.com/2019/07/08/keras-imagedatagenerator-and-data-augmentation/>)

Glossary:

- **Activation**: The probability of a neuron being in the on-state.
- **Classifier**: Categorizing an image into a class.
- **Convolution**: matrix multiplication of a kernel with the pixel value to get a desired effect.
- **Dataset**: The input features on which the neural network has to be trained.
- **Gradient**: The direction along maximum change.
- **Hyperparameters**: The hard-coded parameters in a neural network. These do not change with time.
- **Overfitting**: The phenomenon where the model is generalized to the training dataset and does not perform satisfactorily in testing.
- **Pooling**: A method where the maximum value in a neighbourhood is chosen to reduce the size of the image.
- **Regression**: A measure of the relation between the mean value of one variable and corresponding values of other variables.
- **Vectorization**: A method where numbers are stacked in the form of a matrix for the ease of computation.